

Contents

Connect Four — Implementation Plan

Status at a glance (2026-05-17)

Goal

Design principle: architecture before game

Phase A — Architectural alignment (TTT only, no C4 code)

Phase B — Connect 4 implementation

Phase C — Polish

Risks + mitigations

Sequencing summary

Open items (will be answered as the plan executes)

1

1

2

2

3

6

7

8

9

9

Connect Four — Implementation Plan

Status: Draft, 2026-05-17. Supersedes Platform_Implementation_Plan.md §Phase 4 — the original “drop in another game package” framing pre-dated the 2026 design lock that elevated Connect 4 into a multi-phase architectural initiative.

Related docs: - Game_SDK_Developer_Guide.md — contract every game package must conform to - Platform_Implementation_Plan.md — broader platform roadmap; Phase 3.8 (multi-skill bots) is the prerequisite and is shipped - Help_Corpus/matches-and-games.md, match-formats.md, match-design-rationale.md, solved-games-and-master.md — match semantics + Master tier (already published to corpus) - Help_Corpus/algorithm-history-*.md, the-history-of-game-ai.md — algorithm background pack (already published to corpus)

Status at a glance (2026-05-17)

Sprint	Status	Items	Notes
Phase A — Architectural alignment (TTT only)			
A1 — SDK + game-as-prefix routing + slug rename	shipped	10/10	Legacy xo alias kept active via LEGACY_SLUG_MAP; 301 redirects + map removal moved to C6. Regression sweep + staging smoke green on v1.4.0-alpha-5.11.
A2 — Match-based play + match-level ELO	shipped	11/11	Historical rows frozen at cutover; new matches use the new formula. Corpus alignment audit landed an exact worked-example test pinned to match-formats.md.
A3a — Training data + UX rework (in-process)	in progress	5/15	Presets (Quick/Standard/Deep) per (game, algorithm) + ETA computed up front and persisted onto the session. Earlier: schema audit, channel rename, tournament ELO + series-winner attribution.
A3b — Training worker process cutover	not started	0/9	New xo-training Fly app, Redis queue, pub/sub streaming.

Sprint	Status	Items	Notes
A4 — Multi-skill bot UI + auto-clone	not started	0/13	Phase 3.8 data layer is already shipped; this finishes the UI.
Phase B — Connect 4 build			
B1 — packages/game-connect-four/package	blocked on A1+A3	0/7	Engine, Master solver, no UI.
B2 — C4 play surface (desktop + mobile)	blocked on B1	0/9	6×7 board, mobile column-input (tap-to-preview, tap-to-confirm).
B3 — C4 ranked + tournament	blocked on A2+B2	0/5	Reuses match infra from A2; no new ELO code.
B4 — C4 training, all five algorithms	blocked on A3+B1	0/8	Rule-Based / Minimax / MCTS / DQN / AlphaZero.
B5 — Master tier + Master Challenge	blocked on B1+B2	0/11	Perfect-play solver, off-ladder, best-of-2 Challenge format.
Phase C — Polish			
C1 — Journey + coaching updates	not started	0/3	New onboarding milestone post-first-ranked-TTT-win.
C2 — Training UX completions	not started	0/3	Stop-early, 7-day rollback, advanced disclosure.
C3 — Corpus completions	not started	0/3	Cost/preset docs queue + C4-specific examples.
C4 — Accessibility + mobile polish	not started	0/3	Keyboard nav, screen-reader labels, gesture conflicts.
C5 — Observability sweep	not started	0/2	C4 dashboards + alert thresholds.
C6 — Legacy slug cleanup	blocked on B exit	0/3	Drop LEGACY_SLUG_MAP, add 301 redirects from /xo*, scrub residual xo literals.
Total		21/114	

Update this table as sprints land. Each – [] flipped to – [x] in the body should be reflected in the `Items` column.

Goal

Ship Connect 4 as the platform's second game, alongside Tic-Tac-Toe, with:

- Strict SDK conformance (game logic lives entirely in `packages/game-connect-four/`; if the SDK is deficient, extend the SDK rather than smuggling game logic into the platform).
- Match-based play with alternating colors and match-level ELO (best-of-2 ranked, best-of-3 tournament, casual stays best-of-1).
- All five training algorithms (Rule-Based, Minimax, MCTS, DQN, AlphaZero) operating through a unified async training architecture with realtime W/D/L streaming, multi-curve eval, preset gating, checkpoint resume, and a dedicated worker process.
- A Master tier (perfect-play minimax solver, off-ladder) for the AlphaZero-teaching moment.
- Mobile column-input UI (tap-to-preview, tap-to-confirm).

Design principle: architecture before game

The first half of the plan touches **no Connect 4 code at all**. Every architectural change required by Connect 4 (slug rename, game-as-prefix routing, match-based play, training rework, worker process, multi-skill bot UI) is implemented and validated

against the existing Tic-Tac-Toe game first.

Rationale: each architectural change is a potential regression vector. Doing them under the cover of “we’re shipping a new game” tangles new-feature risk with refactor risk. Doing them on TTT first lets us regress-detect against a familiar surface, with no rollback ambiguity. Only after the platform reaches architectural parity does Connect 4 code begin.

This is enforced in the phase structure below: **Phase A is TTT-only, Phase B is the C4 build, Phase C is platform polish.**

Phase A — Architectural alignment (TTT only, no C4 code)

Five sprints. Each ships through the normal dev → staging → main deploy flow. End state: TTT works identically to today from a user’s perspective, but the platform underneath is ready to host a second game.

A1 — SDK audit + game-as-prefix routing + slug rename

Goal: Move TTT onto the routes and slug that the platform will use long-term, with the SDK contract extended to cover anything Connect 4 will need.

- ☒ SDK audit — confirm the current contract (per `Game_SDK_Developer_Guide.md`) suffices for a 6×7 column-input game with a perfect-play “Master” tier. Identify any gaps.
- ☒ Extend SDK if needed (likely: `meta.inputMode: 'cell' | 'column'`, `meta.matchFormat` hints, `meta.masterStrategy` hook for off-ladder perfect-play bots). Bump SDK version; update `Game_SDK_Developer_Guide.md`.
- ☒ Refactor `packages/game-xo/` to consume any new SDK fields explicitly. No hidden TTT assumptions.
- ☒ Add `/games/<slug>/...` route layer to backend (`/api/v1/games/<slug>/...`), landing (`/games/<slug>/...`), and tournament service.
- ☒ DB migration: rename `xo` → `tic-tac-toe` in `BotSkill.gameId`, `GameElo.gameId`, `Tournament.game`, and any other `gameId`-bearing rows. Single migration; no dual-write phase.
- ☒ Backend code: strip `GAME_ID = 'xo'` constant from `eloService.js`; parameterize every `gameId`-aware function.
- ☒ Backend cache keys: rename `bots:gameId:xo` → `bots:gameId:tic-tac-toe`. Flush on deploy.
- ☒ Landing: update `BotFilterBar.GAMES`, route definitions, deep-link parsers, picker UI.
- ☒ Update all docs in `/doc/Help_Corpus/` that reference the slug `xo` (search-and-replace pass).
- ☒ Tests: regression suite passes; e2e journey + smoke pass on staging.

Moved to C6 — 301 redirects from `/xo*` paths and `LEGACY_SLUG_MAP` removal. Doing them here would shorten the deprecation window to days; doing them in C6 (after Phase B ships) gives external callers — cached client bundles, bookmarks, anyone polling our API — multiple release cycles to migrate.

A2 — Match-based play + match-level ELO (TTT)

Goal: Ranked TTT becomes best-of-2 with alternating colors. Tournament becomes best-of-3 with random-color game-3 tiebreaker. ELO updates per match, not per game. Casual stays best-of-1.

- ☒ New Match entity in DB (or extend Game with `matchId`, `matchSequence`). Decide schema: pure relational Match row with child Game rows, or denormalized `matchId` column on Game.
- ☒ Match lifecycle states: `FORMING` → `IN_PROGRESS` → `COMPLETED`, with per-game results aggregated to a match score.
- ☒ Color assignment logic — game 1 random/seeded, subsequent games swap, game 3 tournament tiebreaker random with announcement.
- ☒ Ranked play: server orchestrates a best-of-2 match for any “Ranked vs X” entry point. No client-side color choice.
- ☒ Tournament play: extend existing `TournamentMatch` (`p1Wins`, `p2Wins`, `drawGames` already exist) to enforce best-of-3 with the random-color rule for game 3.
- ☒ `eloService.js`: implement match-score ELO update (sum of per-game 1.0/0.5/0.0 scores ÷ N games), replacing the per-game update for ranked/tournament. Casual single-game update unchanged.
- ☒ Historical rows: stay in place. No retroactive recompute. Pre-cutover ELO is the starting point; new matches move from there.
- ☒ Realtime: emit match-level events (`match.started`, `match.gameComplete`, `match.completed`) on the existing realtime channel. Client renders match progress (1-0, 1-1, etc.).
- ☒ Landing UI: ranked entry points say “Best of 2” explicitly. Match progress visible during play.
- ☒ Tests: ELO math worked-example test (matches the example in `Help_Corpus/match-formats.md`). E2E ranked-match flow. Tournament best-of-3 with 1-1 tiebreaker. (*Tournament BO3 e2e remains to add when tournament ELO moves to*

per-match — see corpus footnote.)

- ✗ Corpus alignment: matches-and-games.md, match-formats.md, match-design-rationale.md are already published — confirm they describe the shipped behavior. (A2.8 audit: fixed the “smaller weight” casual claim, replaced approximate worked-example numbers with the exact +1.8 / +17.8 / -14.2 deltas the formula produces, footnoted that tournament BO3 still updates ELO per-game today.)

A3a — Training data + UX rework (in-process, TTT)

Goal: Reshape the training API into the form Connect 4 will need, without yet moving training out of the backend process. TTT validates everything.

- ✗ **Tournament series-winner attribution: fix mark-vs-participant aggregation.** Surfaced during the A3a.1 audit. socketHandler.recordPvpGame’s tournament branch was persisting TournamentMatch.p1Wins = xWins and p2Wins = oWins and computing the series winner as xWins >= oWins ? 'X' : 'O' — both wrong whenever colors swap, which rematchGame’s tournament branch always does (game 1: seat1=X, game 2: seat1=O, game 3: random). Scenarios like *seat1 wins game 1 as X, seat1 wins game 2 as O* were getting stored as xWins=1, oWins=1 (a “tie”) when seat1 had actually swept 2-0. Fix: pulled the tournament branch’s aggregation up to read Game.winnerId per tournamentMatchId (participant-keyed) and resolve host/guest → TournamentMatch.participant1Id / participant2Id by querying the row; both the series-done complete call and the mid-series persist now write participant-keyed totals. Tied series default to host (preserves the pre-A3a.1.5 “tied → X” fallback now that we no longer assume X=host across the whole series). Tests: 8 total — flag-on/off ELO, draw handling, HvB exclusion, 2-0 sweep with color-swap mismatch (the headline bug), inverted host/participant orientation, mid-series persist, all-draw BO3 fallback to host.
- ✗ **A2 carry-over: tournament match-level ELO (opt-in per game).** Audit during A3a kickoff turned up two facts the previous footnote got wrong: (a) tournament matches never updated ELO at all (the casual PvP ELO call has been guarded by !isTournamentRoom since tournament play first landed in 3843f49 — “ELO update: skip for tournament games (design requirement)”), and (b) the X/O mark-aggregated xWins/oWins the tournament branch carries aren’t participant-relative because colors swap game-to-game. The fix is intentionally opt-in: a tournamentMovesElo flag on each game’s SDK meta, mirrored into a TOURNAMENT_ELO_GAME_IDS set on the backend (backend/src/constants/games.js). TTT opts out (BO3 game-3 is a coinflip on a solved game); Connect 4 will opt in at launch. When the flag is on, socketHandler.recordPvpGame reads Game rows by tournamentMatchId and aggregates winnerId (participant-keyed) into a single updateBothElosAfterMatch call. Unit tests cover both flag states + draw handling + HvB exclusion. Corpus footnote in match-formats.md rewritten to describe the actual rule.
- ✗ **TrainingSession audit + TrainingCheckpoint + TrainingMetric.** Existing TrainingSession covered modelId/mode/iterations/status/config/summary/timestamps + TrainingEpisode children, but was missing presets, ETA gating fields, approval audit trail, pause/resume timestamps, and any checkpoint pointer. Additive migration (20260519020000_a3a_training_session_audit) added: preset, expectedDurationMs, approvalStatus, approvalRequestedAt, approvedAt, approvedById (FK → User), pausedAt, checkpointEpisode — all nullable so existing rows continue to work. New training_checkpoints table (sessionId, episodeNum, weights, runtimeState; unique on (sessionId, episodeNum)) supports checkpoint-based recovery. New training_metrics table (sessionId, episodeNum, opponentLabel, wins/draws/losses, asFirstMover?; indexed on (sessionId, opponentLabel, episodeNum)) supports stacked W/D/L curves + the deferred Master split. 6 schema tests cover field write/read, unique-checkpoint enforcement, multi-curve metric writes, asFirstMover, and FK cascade.
- ✗ **Channel rename: ml:session:{id} → training:<sessionId>.** One-for-one rename across 7 files (backend services + routes + landing UI + tests): channel prefix ml:session: → training: and each event type ml:progress / ml:complete / ml:cancelled / ml:error / ml:curriculum_advance / ml:early_stop rebadged to training:*. The existing “doubled-up” topic shape (training:abc:training:progress) is preserved — collapsing that to training:abc:progress would touch more files and isn’t required for the SDK contract. Backend 2129/2129, landing 507/507 still green after the cut.
- ✗ **Preset model: Quick / Standard / Deep per (game, algorithm).** Catalog lives in backend/src/config/trainingPresets.js: one entry per (gameId × algorithm × preset) returning { iterations, expectedDurationMs }. TTT covers all 7 trainable algorithms (qlearning, sarsa, montecarlo, policygradient, dqn, alphazero, mcts) + rule_based (zeroed presets, labelling-only). resolvePreset() normalises algorithm casing (Q_LEARNING / q-learning / qlearning all resolve to the same entry). Wired into mlService.startTraining + startFrontendSession and the equivalent skillService functions: preset overrides the iterations arg, persists preset + expectedDurationMs on the session row, throws “Unknown preset” for bad combinations (route returns 400). Catalog deliberately puts AlphaZero standard/deep + DQN deep + policygradient deep past the 4-hour mark so A3a.5 gating gets exercised on every trainable algorithm. Per-session iteration cap is preserved — “deep” qlearning/sarsa/montecarlo land at 100k, the hard cap. Tests: 12 catalog/resolver tests + 7 startTraining wiring tests.

- ☐ ETA-threshold gating — sessions ≥ 4 hours require admin approval. Sub-4hr sessions auto-confirm. Approval flow has a queue UI for admins.
- ☐ Per-user concurrency cap: 1 active training session per user (configurable per role).
- ☐ Multi-curve eval: each session periodically plays N games against the algorithm-appropriate primary opponent + Easy + Medium side curves. Budget-split sampling (~12 primary + 4 Easy + 4 Medium, total ~20 games per eval), keeping overhead ~5%.
- ☐ Stacked W/D/L visualization: each eval point emits W/D/L tallies. Frontend renders stacked area chart per opponent curve. Saturated curves (e.g., 100% wins against Easy) auto-fade.
- ☐ Master-vs-bot split (deferred wiring; lights up in Phase B): if eval includes Master, split into “as first-mover” and “as second-mover” sub-charts. No-op for TTT (no Master tier).
- ☐ Checkpoint-based recovery: every N episodes, persist a checkpoint. On process crash, session resumes from the last checkpoint on next backend boot.
- ☐ Resume button: users can pause and resume sessions (not yet “stop early”). Stop-early is a Phase C item.
- ☐ No knobs exposed in the v1 UI — presets only. Advanced disclosure deferred.
- ☐ Tests: TrainingSession + TrainingMetric CRUD, preset selection, ETA gating boundary at 4hr, multi-curve eval budget math, checkpoint resume after simulated crash.
- ☐ Corpus: write the queued training UX docs (how sessions work, reading W/D/L charts, reading multi-curve view, when to stop, time/cost, resuming, self-play vs eval).

A3b — Training worker process cutover (TTT)

Goal: Move training out of the backend process into a dedicated worker, so a long AlphaZero Deep run can never starve the API box.

- ☐ New Fly app: xo-training-prod (and -staging), sharing the backend container image. Separate scale + resources from the API backend.
- ☐ Redis-backed job queue (BullMQ or equivalent; we already have Redis in the stack). Backend enqueues startTraining jobs; worker(s) consume.
- ☐ Pub/sub channel: worker emits training events to Redis; backend SSE handler subscribes and fans out to connected clients. Realtime contract from A3a is unchanged (clients still consume training:<sessionId> via the backend).
- ☐ Graceful shutdown: worker handles SIGTERM by checkpointing the current session, persisting state, and exiting cleanly. New worker on next deploy resumes from checkpoint.
- ☐ Cap enforcement at the queue: per-user cap (1 active) + global cap (configurable, e.g., 4 concurrent sessions) gated before a job leaves the queue.
- ☐ Failure handling: job retries on transient errors, dead-letter after N failures. Sessions in dead-letter surface to admin with a “diagnostic dump” link.
- ☐ Observability: per-worker CPU/RAM dashboard, queue depth alert, dead-letter alert. Hook into Observability_Plan.md.
- ☐ Tests: queue enqueue/consume happy path, cap enforcement, graceful-shutdown checkpoint, crash-during-training recovery, dead-letter flow.
- ☐ Deploy: fly launch for xo-training-staging, wire to staging Redis/DB. Validate TTT training under load (multiple concurrent users, mix of algorithms). Promote to prod when green.

A4 — Multi-skill bot UI + auto-clone

Goal: Finish the Phase 3.8 remainder. By the time Connect 4 ships, multi-game bot management is already in users’ hands.

- ☐ Bot detail page: per-skill cards listing every BotSkill row owned by the bot, with per-game ELO and last-trained timestamps.
- ☐ Add-skill flow: from a bot, pick a game the bot doesn’t have a skill for, pick an algorithm, optionally clone hyperparams from an existing skill. Wire to repointBotPrimarySkill so User.botModelId updates correctly.
- ☐ Auto-clone non-primary skills on rename/rebrand: when the bot is renamed, every skill carries along (already true; verify and test).
- ☐ Picker disambiguation: when a bot has multiple skills, the picker shows “BotName (game)” with the active gameId, so users always know what they’re queuing.
- ☐ Profile + Gym nav: deep-link straight to a skill’s training page from the bot card.
- ☐ Pre-C4 validation: with only TTT live, the UI is intentionally lightly populated. We’re validating the data plumbing, not showcasing breadth.

- Tests: add-skill happy path, rename across multiple skills, picker disambiguation, primary-skill repoint after training completes.

Phase A exit criteria

- TTT plays identically end-to-end (casual + ranked + tournament) on the new routes and new slug.
 - TTT match-level ELO produces sensible movements for the existing user base.
 - All 5 TTT algorithms train through the worker process with W/D/L streaming, multi-curve eval, and checkpoint resume.
 - Multi-skill bot UI lets a user add, manage, and queue from multiple skills (limited to one game today).
 - Help corpus updated: any TTT-facing docs that reference the old slug, single-game ranked, or the old training UX are refreshed.
 - No open regressions on the V1 acceptance script.
-

Phase B — Connect 4 implementation

Five sprints. Each builds atop the architecture established in Phase A. If any sprint here surfaces an SDK gap, the work pauses, the SDK is extended (with `Game_SDK_Developer_Guide.md` updated), and the sprint resumes — game logic never leaks into the platform.

B1 — `packages/game-connect-four/` package

Goal: A standalone, SDK-conformant game package. Engine + bot interface + tests, no UI yet.

- Package scaffold: `packages/game-connect-four/` with the same shape as `packages/game-xo/`.
- Engine: 6×7 board, gravity mechanic, win detection (horizontal/vertical/both diagonals), draw detection (board-full), legal-move enumeration (top-of-column).
- State serialization: compact format suitable for replay storage and bot weight inputs.
- meta export: `id: 'connect-four', inputMode: 'column', matchFormat: { ranked: 'bo2', tournament: 'bo3', master: 'bo2' }, masterStrategy: 'solver', builtInBots: [easy, medium, hard, master]`.
- botInterface:
 - Built-in minimax tiers (Easy depth-2, Medium depth-4, Hard depth-6 to depth-8, Master = the solver).
 - Training hooks for the five algorithms (delegating to the shared `@xo-arena/ai` engines where game-agnostic; new code only where C4 specifics matter).
- Master solver: perfect-play minimax with full-depth search and the known solved-game heuristics (center-first opening, threat parity). Off-ladder.
- Unit tests: engine correctness (win/draw/illegal-move), serialization round-trip, minimax depth correctness, solver vs known opening lines.

B2 — C4 play surface (desktop + mobile)

Goal: Casual quick match against built-in Easy/Medium/Hard works end-to-end via `/games/connect-four/...`

- Game component: 6×7 board with Red/Yellow tokens, drop animation, win highlight, draw indicator.
- Desktop input: click on column to drop.
- Mobile input: tap-to-preview (semi-transparent token at the top of the column showing where the piece will land), tap-to-confirm (drops the piece). Tap a different column to switch preview. Long-press cancels.
- Audio cues: drop, win, draw — reusing the existing `sdk.playSound()` pattern with new sample IDs.
- Theming: Red/Yellow color scheme conforms to platform tokens (`packages/sdk` theme exports).
- Spectate: rendered-table mode works; piece drops animate for spectators.
- Replay: replays of C4 games render via the same replay infrastructure as TTT.
- Bot vs human: server-side bot moves work for built-in Easy/Medium/Hard. Master is not yet wired here (Phase B5).
- Tests: e2e casual match vs each built-in tier on desktop + mobile viewports.

B3 — C4 ranked + tournament

Goal: Ranked best-of-2 and tournament best-of-3 work for C4 with zero new infrastructure — A2's match layer just lights up.

- Wire C4 into ranked match orchestration. Alternating colors, ELO updates per match.

- ☐ Wire C4 into the tournament service. Best-of-3 with random-color game 3 tiebreaker.
- ☐ Classification ladder: C4 ELO tracked separately from TTT (already supported by GameElo (userId, gameId) unique constraint).
- ☐ Landing nav: C4 appears in the games picker; BotFilterBar shows both games.
- ☐ Tests: ranked best-of-2 happy path, tournament with 1-1 game-3 tiebreaker, classification ladder shows independent TTT and C4 ratings for the same user.

B4 — C4 training, all five algorithms

Goal: Every algorithm trains through the A3 architecture with no platform-side changes.

- ☐ Rule-Based: heuristic rules (center preference, win-on-move, block-on-move, fork detection, edge play). Zero-time “training.”
- ☐ Minimax (Quick Bots): tier-label bump only, same pattern as TTT Quick Bots — user:<id>:minimax:<tier> for C4.
- ☐ MCTS: classical UCB1 with rollouts. Tunes simulation count, exploration constant, optional rollout policy.
- ☐ DQN: experience replay + target network. Standard preset hyperparams.
- ☐ AlphaZero: policy/value network with PUCT-guided MCTS. Three presets:
 - Quick — ~30 minutes, small network, low simulation count. Auto-confirmed.
 - Standard — ~4 hours, medium network. **Hits admin-approval threshold.**
 - Deep — ~40 hours, full network + high sim count. Always admin-approved.
- ☐ Eval curves: primary opponent per algorithm (e.g., AZ evals against Hard minimax). Easy + Medium side curves render saturated and auto-fade.
- ☐ Streaming UX: TTT users seeing W/D/L curves in realtime today will see the same UI for C4 with no relearning.
- ☐ Tests: each algorithm reaches expected milestones from a known seed in CI-bounded episode counts.

B5 — Master tier + Master Challenge

Goal: Perfect-play Master, off-ladder, with the Master Challenge user flow.

- ☐ Master button on the C4 home: opens a best-of-2 Master Challenge (one game as first-mover, one as Master first-mover).
- ☐ Master plays the perfect-play solver from B1. Move latency budgeted (solver is fast enough at depth-required for solved-game optimality, but cache opening lines).
- ☐ No ELO update on Master matches. Per-bot “vs Master” stat block on the bot detail page (wins / draws / losses).
- ☐ Training integration: AlphaZero bots can opt into “vs Master” eval as an additional curve (showing first-mover-draw rate over time — the AlphaZero teaching moment).
- ☐ Landing UI: Master section on C4 home page explains the off-ladder rule and what “beating Master” means (linking to Help_Corpus/solved-games-and-master.md).
- ☐ Tests: Master Challenge happy path, off-ladder ELO behavior (no rating movement), AZ training with Master-vs-bot eval curve.

Phase B exit criteria

- ☐ C4 ships through /games/connect-four/... for casual, ranked, tournament, and Master Challenge.
- ☐ All five algorithms train through the A3/A3b architecture. AZ Deep on C4 succeeds at least once on staging without affecting API latency.
- ☐ Multi-skill bots can hold both a TTT and a C4 skill, each with independent ELO.
- ☐ Replay, spectate, ranked classification all work for C4 at parity with TTT.
- ☐ Corpus matches reality — the Connect 4 pages in Help_Corpus/ reflect what shipped.

Phase C — Polish

Approximately 3-4 sprints. Lower urgency than Phase B but completes the end-to-end experience.

C1 — Journey + coaching updates

- ☐ Update the onboarding journey to surface a Connect 4 milestone after the user’s first ranked TTT win.
- ☐ Add coaching cards for C4 specifics (column thinking, first-mover advantage, threat parity).
- ☐ Update Intelligent_Guide_Requirements.md and Intelligent_Guide_Implementation_Plan.md.

C2 — Training UX completions

- ☐ Stop-early button on active sessions (with confirmation + “save current checkpoint”).
- ☐ 7-day rollback button on bot detail page: revert a skill’s weights to a previous checkpoint within a 7-day retention window.
- ☐ Advanced disclosure (v1.1): expose hyperparameter knobs behind an “Advanced” toggle for users who want to tune past presets.

C3 — Corpus completions

- ☐ Write the queued cost/preset docs: “What each preset does”, “How long does training take?”, “What does training cost?”, “What happens if I close the tab?”, “Stop training early”, “Why no knob tweaking”, “Why AZ slow on C4 fast on TTT”.
- ☐ Refresh bots-overview.md, gym-train-tab.md, understanding-training-charts.md with C4-specific examples.
- ☐ Render PDFs via the tuned pandoc invocation.

C4 — Accessibility + mobile polish

- ☐ Keyboard navigation for the C4 board (arrow keys + enter to drop).
- ☐ Screen-reader labels for board state (“Column 4, row 3, your turn”).
- ☐ Mobile UX refinement pass — confirm the tap-to-preview-tap-to-confirm gesture works cleanly across viewports and doesn’t conflict with browser swipe gestures.

C5 — Observability sweep

- ☐ C4-specific dashboards: completion rate per tier, AZ Master-draw-rate over time, training queue depth, worker CPU/RAM under typical load.
- ☐ Alert thresholds: dead-letter queue depth, training session failures per hour, API latency regression during heavy training.

C6 — Legacy slug cleanup

Goal: Drop the `xo` → `tic-tac-toe` alias once Phase B has shipped and external callers have had multiple release cycles to migrate. Cleanup, not architecture — sequenced here on purpose so the deprecation window is long enough that nothing in the wild still calls the old slug.

Pre-conditions: Phase B is shipped to production. No staging or production access logs show requests with `gameId=xo` or path `/play/xo*` for at least one release cycle (≥ 2 weeks). Verify via the request logs / observability dashboards before starting.

- ☐ Add 301 Moved Permanently redirects in `landing/server.js` and (if needed) the backend for legacy paths: `/play/xo*` → `/games/tic-tac-toe/play*`, `/xo/*` → `/games/tic-tac-toe/*`. Keep the redirects in place for one further release.
- ☐ Remove `LEGACY_SLUG_MAP` from `backend/src/constants/games.js` and `tournament/src/constants/games.js`. `resolveGameSlug('xo')` should now return null (caller’s ?? `rawGameId` fallback then surfaces “Unknown game slug” 404 via `validateGameSlug`).
- ☐ Sweep for any residual `'xo'` literals in code (excluding migration files, historical perf-trace doc citations, and `Help_Corpus` search-alias tags). Update or delete each one.
- ☐ Update `Game_SDK_Developer_Guide.md` to drop the “legacy alias accepted” caveat — kebab-case canonical slugs only.

Risks + mitigations

Risk	Mitigation
Slug-rename migration corrupts production data. Renaming <code>xo</code> → <code>tic-tac-toe</code> across <code>BotSkill</code> , <code>GameElo</code> , <code>Tournament</code> is a write-heavy migration.	Dry-run on staging with a prod data snapshot. Migration runs inside a transaction with explicit row counts logged. Rollback plan: a reverse-rename migration is staged before deploy.
Match-level ELO produces unexpected user-visible drops or jumps at cutover. Even with frozen historical rows, the first few matches under the new formula will move ratings under different math.	Communicate proactively in release notes. The new math gives smaller per-match deltas (it’s “match score 0.5” vs “game-loss 0.0”), so the practical effect is gentler movement. Monitor for outliers.

Risk	Mitigation
AlphaZero Deep run on C4 starves the API box despite the worker. Per-user cap and worker isolation should prevent this, but a misconfigured Fly app could fall back to in-process.	Pre-launch load test on staging: queue a real AZ Deep run while a separate user plays ranked. API p95 latency must not regress.
SDK gap surfaces mid-Phase B. Discovering during B2 that the SDK can't express column-input semantics would block the sprint.	Phase A1 audits explicitly for this. If a gap appears in B, the SDK is extended in place; we don't ship workarounds.
Master solver is too slow at depth-required. Perfect-play C4 needs deep search; if move latency is too high, the Master tier feels broken.	Cache the entire opening book (first 10-12 plies) as a lookup table. Worst case, fall back to a precomputed move table for the opening, switch to live search after the book exits.
Mobile column-input gesture conflicts with browser scroll. A column tap shouldn't scroll the page.	touch-action: manipulation on the board; e2e tests on mobile viewports.
Worker deployment introduces a new failure surface. A worker outage could halt all training.	Status page reflects worker health. If the worker is down, new training-start requests return a clear "training temporarily unavailable" error and queue UI surfaces the outage. Existing ranked/casual/tournament play is unaffected (those don't go through the worker).

Sequencing summary

Phase A (TTT only):

- A1 – SDK + routing + slug rename
- A2 – Match-based play + match-level ELO
- A3a – Training data + UX rework (in-process)
- A3b – Worker process cutover
- A4 – Multi-skill bot UI + auto-clone

Phase B (Connect 4 build):

- B1 – packages/game-connect-four/
- B2 – C4 play surface (desktop + mobile)
- B3 – C4 ranked + tournament
- B4 – C4 training, all five algorithms
- B5 – Master tier + Master Challenge

Phase C (polish):

- C1 – Journey + coaching
- C2 – Training UX completions
- C3 – Corpus completions
- C4 – Accessibility + mobile polish
- C5 – Observability sweep

Phases ship sequentially. Sprints within a phase can run in parallel where they don't have ordering dependencies (e.g., B2 and B3 can overlap; B4 needs the C4 package from B1; B5 needs the C4 surface from B2).

Open items (will be answered as the plan executes)

- **Worker pool sizing.** Start with count=1 on staging, count=2 on prod, scale based on observed queue depth. Final sizing TBD post-launch.
- **AZ C4 Deep economics.** Per the corpus, AZ C4 Deep is roughly \$3–50 per run depending on Fly machine class. The exact admin-approval UX (credits charged? burned against an admin pool?) is open — pin down before Phase B4.
- **Master solver move-cache size.** Opening-book lookup table size depends on how deep the book runs. Spike during B1.
- **Color naming.** Red/Yellow is canonical; settle whether colorblind-mode swaps to a high-contrast pair or uses pattern fills. Decide during B2.