

Contents

Connect Four — Implementation Plan

Status at a glance (2026-05-17)

Goal

Design principle: architecture before game

Phase A — Architectural alignment (TTT only, no C4 code)

Phase B — Connect 4 implementation

Phase C — Polish

Risks + mitigations

Sequencing summary

Open items (will be answered as the plan executes)

1

1

2

3

3

7

9

10

11

11

Connect Four — Implementation Plan

Status: Draft, 2026-05-17. Supersedes Platform_Implementation_Plan.md §Phase 4 — the original “drop in another game package” framing pre-dated the 2026 design lock that elevated Connect 4 into a multi-phase architectural initiative.

Related docs: - Game_SDK_Developer_Guide.md — contract every game package must conform to - Platform_Implementation_Plan.md — broader platform roadmap; Phase 3.8 (multi-skill bots) is the prerequisite and is shipped - Help_Corpus/matches-and-games.md, match-formats.md, match-design-rationale.md, solved-games-and-master.md — match semantics + Master tier (already published to corpus) - Help_Corpus/algorithm-history-*.md, the-history-of-game-ai.md — algorithm background pack (already published to corpus)

Status at a glance (2026-05-17)

Sprint	Status	Items	Notes
Phase A — Architectural alignment (TTT only)			
A1 — SDK + game-as-prefix routing + slug rename	shipped	10/10	Legacy xo alias kept active via LEGACY_SLUG_MAP; 301 redirects + map removal moved to C6. Regression sweep + staging smoke green on v1.4.0-alpha-5.11.
A2 — Match-based play + match-level ELO	shipped	11/11	Historical rows frozen at cutover; new matches use the new formula. Corpus alignment audit landed an exact worked-example test pinned to match-formats.md.
A3a — Training data + UX rework (in-process)	shipped	12/12	All backend + corpus + tests done. Three UI items (stacked W/D/L viz, Master-vs-bot split, “no knobs” v1 disclosure) moved to A3b where they ship alongside the worker cutover.

Sprint	Status	Items	Notes
A3b — Training worker process cutover + training UI	in progress	1/10	Sequenced as 10 PR-sized chunks. Shipped: A3b.1 spike — BullMQ + new xo-training worker round-trip working end-to-end locally; Dockerfile fix so workspace installs pick up bullmq's transitives. Next: A3b.2a (worker consumes <code>training:start</code> in shadow mode).
A4 — Multi-skill bot UI + auto-clone	not started	0/13	Phase 3.8 data layer is already shipped; this finishes the UI.
Phase B — Connect 4 build			
B1 — packages/game-connect-four/package	blocked on A1+A3	0/7	Engine, Master solver, no UI.
B2 — C4 play surface (desktop + mobile)	blocked on B1	0/9	6×7 board, mobile column-input (tap-to-preview, tap-to-confirm).
B3 — C4 ranked + tournament	blocked on A2+B2	0/5	Reuses match infra from A2; no new ELO code.
B4 — C4 training, all five algorithms	blocked on A3+B1	0/8	Rule-Based / Minimax / MCTS / DQN / AlphaZero.
B5 — Master tier + Master Challenge	blocked on B1+B2	0/11	Perfect-play solver, off-ladder, best-of-2 Challenge format.
Phase C — Polish			
C1 — Journey + coaching updates	not started	0/3	New onboarding milestone post-first-ranked-TTT-win.
C2 — Training UX completions	not started	0/3	Stop-early, 7-day rollback, advanced disclosure.
C3 — Corpus completions	not started	0/3	Cost/preset docs queue + C4-specific examples.
C4 — Accessibility + mobile polish	not started	0/3	Keyboard nav, screen-reader labels, gesture conflicts.
C5 — Observability sweep	not started	0/2	C4 dashboards + alert thresholds.
C6 — Legacy slug cleanup	blocked on B exit	0/3	Drop <code>LEGACY_SLUG_MAP</code> , add 301 redirects from <code>/xo*</code> , scrub residual xo literals.
Total		21/114	

Update this table as sprints land. Each `- [ ]` flipped to `- [x]` in the body should be reflected in the `Items` column.

Goal

Ship Connect 4 as the platform's second game, alongside Tic-Tac-Toe, with:

- Strict SDK conformance (game logic lives entirely in `packages/game-connect-four/`; if the SDK is deficient, extend the SDK rather than smuggling game logic into the platform).

- Match-based play with alternating colors and match-level ELO (best-of-2 ranked, best-of-3 tournament, casual stays best-of-1).
- All five training algorithms (Rule-Based, Minimax, MCTS, DQN, AlphaZero) operating through a unified async training architecture with realtime W/D/L streaming, multi-curve eval, preset gating, checkpoint resume, and a dedicated worker process.
- A Master tier (perfect-play minimax solver, off-ladder) for the AlphaZero-teaching moment.
- Mobile column-input UI (tap-to-preview, tap-to-confirm).

Design principle: architecture before game

The first half of the plan touches **no Connect 4 code at all**. Every architectural change required by Connect 4 (slug rename, game-as-prefix routing, match-based play, training rework, worker process, multi-skill bot UI) is implemented and validated against the existing Tic-Tac-Toe game first.

Rationale: each architectural change is a potential regression vector. Doing them under the cover of “we’re shipping a new game” tangles new-feature risk with refactor risk. Doing them on TTT first lets us regress-detect against a familiar surface, with no rollback ambiguity. Only after the platform reaches architectural parity does Connect 4 code begin.

This is enforced in the phase structure below: **Phase A is TTT-only, Phase B is the C4 build, Phase C is platform polish.**

Phase A — Architectural alignment (TTT only, no C4 code)

Five sprints. Each ships through the normal dev → staging → main deploy flow. End state: TTT works identically to today from a user’s perspective, but the platform underneath is ready to host a second game.

A1 — SDK audit + game-as-prefix routing + slug rename

Goal: Move TTT onto the routes and slug that the platform will use long-term, with the SDK contract extended to cover anything Connect 4 will need.

- ☒ SDK audit — confirm the current contract (per `Game_SDK_Developer_Guide.md`) suffices for a 6×7 column-input game with a perfect-play “Master” tier. Identify any gaps.
- ☒ Extend SDK if needed (likely: `meta.inputMode: 'cell' | 'column'`, `meta.matchFormat` hints, `meta.masterStrategy` hook for off-ladder perfect-play bots). Bump SDK version; update `Game_SDK_Developer_Guide.md`.
- ☒ Refactor `packages/game-xo/` to consume any new SDK fields explicitly. No hidden TTT assumptions.
- ☒ Add `/games/<slug>/...` route layer to backend (`/api/v1/games/<slug>/...`), landing (`/games/<slug>/...`), and tournament service.
- ☒ DB migration: rename `xo` → `tic-tac-toe` in `BotSkill.gameId`, `GameElo.gameId`, `Tournament.game`, and any other `gameId`-bearing rows. Single migration; no dual-write phase.
- ☒ Backend code: strip `GAME_ID = 'xo'` constant from `eloService.js`; parameterize every `gameId`-aware function.
- ☒ Backend cache keys: rename `bots:gameId:xo` → `bots:gameId:tic-tac-toe`. Flush on deploy.
- ☒ Landing: update `BotFilterBar.GAMES`, route definitions, deep-link parsers, picker UI.
- ☒ Update all docs in `/doc/Help_Corpus/` that reference the slug `xo` (search-and-replace pass).
- ☒ Tests: regression suite passes; e2e journey + smoke pass on staging.

Moved to C6 — 301 redirects from `/xo*` paths and `LEGACY_SLUG_MAP` removal. Doing them here would shorten the deprecation window to days; doing them in C6 (after Phase B ships) gives external callers — cached client bundles, bookmarks, anyone polling our API — multiple release cycles to migrate.

A2 — Match-based play + match-level ELO (TTT)

Goal: Ranked TTT becomes best-of-2 with alternating colors. Tournament becomes best-of-3 with random-color game-3 tiebreaker. ELO updates per match, not per game. Casual stays best-of-1.

- ☒ New Match entity in DB (or extend Game with `matchId`, `matchSequence`). Decide schema: pure relational Match row with child Game rows, or denormalized `matchId` column on Game.
- ☒ Match lifecycle states: `FORMING` → `IN_PROGRESS` → `COMPLETED`, with per-game results aggregated to a match score.
- ☒ Color assignment logic — game 1 random/seeded, subsequent games swap, game 3 tournament tiebreaker random with announcement.

- ☒ Ranked play: server orchestrates a best-of-2 match for any “Ranked vs X” entry point. No client-side color choice.
- ☒ Tournament play: extend existing TournamentMatch (p1Wins, p2Wins, drawGames already exist) to enforce best-of-3 with the random-color rule for game 3.
- ☒ eloService.js: implement match-score ELO update (sum of per-game 1.0/0.5/0.0 scores ÷ N games), replacing the per-game update for ranked/tournament. Casual single-game update unchanged.
- ☒ Historical rows: stay in place. No retroactive recompute. Pre-cutover ELO is the starting point; new matches move from there.
- ☒ Realtime: emit match-level events (match.started, match.gameComplete, match.completed) on the existing realtime channel. Client renders match progress (1-0, 1-1, etc.).
- ☒ Landing UI: ranked entry points say “Best of 2” explicitly. Match progress visible during play.
- ☒ Tests: ELO math worked-example test (matches the example in Help_Corpus/match-formats.md). E2E ranked-match flow. Tournament best-of-3 with 1-1 tiebreaker. *(Tournament BO3 e2e remains to add when tournament ELO moves to per-match — see corpus footnote.)*
- ☒ Corpus alignment: matches-and-games.md, match-formats.md, match-design-rationale.md are already published — confirm they describe the shipped behavior. *(A2.8 audit: fixed the “smaller weight” casual claim, replaced approximate worked-example numbers with the exact +1.8 / +17.8 / -14.2 deltas the formula produces, footnoted that tournament BO3 still updates ELO per-game today.)*

A3a — Training data + UX rework (in-process, TTT)

Goal: Reshape the training API into the form Connect 4 will need, without yet moving training out of the backend process. TTT validates everything.

- ☒ **Tournament series-winner attribution: fix mark-vs-participant aggregation.** Surfaced during the A3a.1 audit. socketHandler.recordPvpGame’s tournament branch was persisting TournamentMatch.p1Wins = xWins and p2Wins = oWins and computing the series winner as xWins >= oWins ? 'X' : 'O' — both wrong whenever colors swap, which rematchGame’s tournament branch always does (game 1: seat1=X, game 2: seat1=O, game 3: random). Scenarios like *seat1 wins game 1 as X, seat1 wins game 2 as O* were getting stored as xWins=1, oWins=1 (a “tie”) when seat1 had actually swept 2-0. Fix: pulled the tournament branch’s aggregation up to read Game.winnerId per tournamentMatchId (participant-keyed) and resolve host/guest → TournamentMatch.participant1Id / participant2Id by querying the row; both the series-done complete call and the mid-series persist now write participant-keyed totals. Tied series default to host (preserves the pre-A3a.1.5 “tied → X” fallback now that we no longer assume X=host across the whole series). Tests: 8 total — flag-on/off ELO, draw handling, HvB exclusion, 2-0 sweep with color-swap mismatch (the headline bug), inverted host/participant orientation, mid-series persist, all-draw BO3 fallback to host.
- ☒ **A2 carry-over: tournament match-level ELO (opt-in per game).** Audit during A3a kickoff turned up two facts the previous footnote got wrong: (a) tournament matches never updated ELO at all (the casual PvP ELO call has been guarded by !isTournamentRoom since tournament play first landed in 3843f49 — “ELO update: skip for tournament games (design requirement)”), and (b) the X/O mark-aggregated xWins/oWins the tournament branch carries aren’t participant-relative because colors swap game-to-game. The fix is intentionally opt-in: a tournamentMovesElo flag on each game’s SDK meta, mirrored into a TOURNAMENT_ELO_GAME_IDS set on the backend (backend/src/constants/games.js). TTT opts out (BO3 game-3 is a coinflip on a solved game); Connect 4 will opt in at launch. When the flag is on, socketHandler.recordPvpGame reads Game rows by tournamentMatchId and aggregates winnerId (participant-keyed) into a single updateBothElosAfterMatch call. Unit tests cover both flag states + draw handling + HvB exclusion. Corpus footnote in match-formats.md rewritten to describe the actual rule.
- ☒ **TrainingSession audit + TrainingCheckpoint + TrainingMetric.** Existing TrainingSession covered modelId/mode/iterations/status/config/summary/timestamps + TrainingEpisode children, but was missing presets, ETA gating fields, approval audit trail, pause/resume timestamps, and any checkpoint pointer. Additive migration (20260519020000_a3a_training_session_audit) added: preset, expectedDurationMs, approvalStatus, approvalRequestedAt, approvedAt, approvedById (FK → User), pausedAt, checkpointEpisode — all nullable so existing rows continue to work. New training_checkpoints table (sessionId, episodeNum, weights, runtimeState; unique on (sessionId, episodeNum)) supports checkpoint-based recovery. New training_metrics table (sessionId, episodeNum, opponentLabel, wins/draws/losses, asFirstMover?; indexed on (sessionId, opponentLabel, episodeNum)) supports stacked W/D/L curves + the deferred Master split. 6 schema tests cover field write/read, unique-checkpoint enforcement, multi-curve metric writes, asFirstMover, and FK cascade.
- ☒ **Channel rename: ml:session:{id} → training:<sessionId>.** One-for-one rename across 7 files (backend services + routes + landing UI + tests): channel prefix ml:session: → training: and each event type ml:progress / ml:complete / ml:cancelled / ml:error / ml:curriculum_advance / ml:early_stop rebadged to training:*. The existing “doubled-up” topic shape (training:abc:training:progress) is preserved — collapsing that to train-

ing:abc:progress would touch more files and isn't required for the SDK contract. Backend 2129/2129, landing 507/507 still green after the cut.

- ☒ **Preset model: Quick / Standard / Deep per (game, algorithm).** Catalog lives in backend/src/config/trainingPresets.js: one entry per (gameId × algorithm × preset) returning { iterations, expectedDurationMs }. TTT covers all 7 trainable algorithms (qlearning, sarsa, montecarlo, policygradient, dqn, alphazero, mcts) + rule_based (zeroed presets, labelling-only). resolvePreset() normalises algorithm casing (Q_LEARNING / q-learning / qlearning all resolve to the same entry). Wired into mlService.startTraining + startFrontendSession and the equivalent skillService functions: preset overrides the iterations arg, persists preset + expectedDurationMs on the session row, throws “Unknown preset” for bad combinations (route returns 400). Catalog deliberately puts AlphaZero standard/deep + DQN deep + policygradient deep past the 4-hour mark so A3a.5 gating gets exercised on every trainable algorithm. Per-session iteration cap is preserved — “deep” qlearning/sarsa/montecarlo land at 100k, the hard cap. Tests: 12 catalog/resolver tests + 7 startTraining wiring tests.
- ☒ **ETA-threshold gating (backend).** In mlService.startTraining / startFrontendSession, after preset resolution: if expectedDurationMs >= ml.approvalThresholdMs (default 4hr, configurable via SystemConfig), the session is created with status='PENDING' + approvalStatus='NEEDED' and the training loop is *not* started. Sub-threshold sessions auto-confirm and run as before. Legacy iterations-direct callers (no expectedDurationMs) are never gated. New service functions approveSession(sessionId, adminUserId) and denySession(sessionId, adminUserId, reason?) enforce state transitions (must be PENDING + NEEDED) and start the loop on approve (queuing if model is busy). Admin endpoints: GET /admin/training/pending, POST /admin/training/:id/approve, POST /admin/training/:id/deny. Tests: 10 new (gate trips for AlphaZero standard, skips for AlphaZero quick, respects SystemConfig override, no-preset never gated, approve idle/busy paths, deny with/without reason, list filter). **Admin queue UI is deferred** to the broader UI sweep (kept out of this sprint per the “no knobs in v1 UI” constraint — backend surface is complete and admin can hit the endpoints directly until then).
- ☒ **Per-user concurrency cap.** _enforceUserConcurrencyCap(ownerBaId) in mlService runs after preset resolution + per-model checks, before both gated and ungated session create. Counts the user's status IN (PENDING, RUNNING) rows joined to BotSkill.createdBy = ownerBaId. Default cap = 1 via ml.maxActiveSessionsPerUser; admins with the ADMIN role get an override via ml.maxActiveSessionsForAdmin (only takes effect when set to a non-zero value, otherwise the default still applies). Skills with createdBy=null (admin-seeded) are skipped — no owner to attribute to. Same check wired into startFrontendSession. Tests: 9 (default cap trip, config override, admin override, non-admin doesn't inherit, no-owner skip, PENDING + RUNNING both count, frontend variant).
- ☒ **Multi-curve eval.** New backend/src/lib/evalCurves.js is a pure module: per-algorithm primary-opponent catalog (ML algorithms → minimax:hard; rule_based → minimax:medium), playEvalGame / runEvalSeries (alternates model's mark across games so first-mover advantage washes out), runEvalBatch (one record per [primary, easy, medium] curve), and recordEvalMetrics (createMany into TrainingMetric). Default budget = 12 primary + 4 easy + 4 medium = 20 games/eval. Wired into _runTraining: every EVAL_GAP=1000 episodes the loop calls runEvalBatch with the engine's no-explore chooseAction(board, false) as the model move fn and minimax of the right tier as the opponent. Persist is fire-and-forget; a transient eval failure logs a warning but never aborts training. 17 module tests (curve resolution, alternating-mark policy, illegal-move forfeit, batch fan-out, budget zero-skip, persistence shape).
- ☒ **Checkpoint-based recovery.** Two-part change. **Persist:** alongside the existing per-model MLCheckpoint write (kept for backward compat), _runTraining now also upserts a per-session TrainingCheckpoint row with { weights, runtimeState: { epsilon } } and updates TrainingSession.checkpointEpisode — both at the existing 1000-episode cadence; upsert means a re-issued checkpoint after resume on the same slot is a no-op. **Resume:** new resumeOrphanedSessions() runs once at backend boot (index.js, fire-and-forget so it never blocks listen()). It finds every status='RUNNING' row, loads the latest TrainingCheckpoint, and either restarts _runTraining with startEpisode + resumedEngineState (weights + epsilon) — or marks the session FAILED and flips BotSkill.status back to IDLE if no checkpoint exists. _runTraining now accepts the resume params: weights replace the model's stored weights, and epsilon overrides the engine's initial value so exploration continues smoothly. Tests: 5 (no-orphans no-op, resume with checkpoint, FAILED no checkpoint, mixed batch, per-session error isolation).
- ☒ **Pause / Resume (backend).** Mirrors the cooperative-cancel pattern. pauseSession(sessionId) adds the session to an in-memory pausedSessions Set and returns the row; the training loop's next tick force-writes a TrainingCheckpoint at the current episode (so resume picks up exactly there, no replay), flushes the in-flight episode batch, transitions the session to PENDING + pausedAt, flips BotSkill.status back to IDLE, emits a training:paused event, and exits the loop without _finishSession so the row stays terminal-free. resumeSession(sessionId) is symmetric: rejects if not paused / no checkpoint, clears pausedAt, flips back to RUNNING, locks the model, and restarts _runTraining with the latest checkpoint's startEpisode + weights + epsilon (or queues if model is busy). resumeOrphanedSessions updated to skip pausedAt-set rows so a boot-time scan doesn't undo a user's explicit pause. Endpoints: POST /api/v1/skills/sessions/:id/pause + .../resume (owner-only via the existing assertModelOwner helper). Stop-

early stays a Phase C concern. UI button deferred. Tests: 9 (pause requires RUNNING + unpaused, resume requires paused + checkpoint, queue when model busy, unknown session, all error paths).

- ☒ **Tests: TrainingSession + TrainingMetric CRUD, preset selection, ETA gating boundary at 4hr, multi-curve eval budget math, checkpoint resume after simulated crash.** Shipped piecewise alongside each item: 6 schema tests (CRUD + checkpoint uniqueness + metric writes + asFirstMover + FK cascade), 12 preset catalog + 7 startTraining wiring, 10 ETA-gate (incl. AlphaZero standard trips, quick skips, SystemConfig override, no-preset never gated, approve/deny paths), 9 per-user cap, 17 multi-curve eval (curve resolution, alternating mark, illegal-move forfeit, batch fan-out, budget zero-skip, persistence shape), 5 checkpoint resume (no-orphans no-op, resume with checkpoint, FAILED no checkpoint, mixed batch, per-session error isolation), 9 pause/resume, plus the um training-recovery harness (10 unit tests on evaluateVerifyState) and the manual 3-session parallel crash-recovery test now run end-to-end against staging.
- ☒ **Corpus.** Four new Help_Corpus docs: **training-presets-and-time** (Quick/Standard/Deep, the 4-hour admin gate, per-user concurrency cap, picking the right preset), **training-multi-curve-eval** (Primary/Easy/Medium reading, saturation fade, what each curve answers, ~20 games/eval budget), **pausing-and-resuming-training** (cooperative pause, crash recovery, pause vs cancel decision), **self-play-vs-eval** (three kinds of games — training episode / eval game / live ranked — and how ELO is updated). Cross-linked back to the existing understanding-training-charts + bot-training-concepts + bot-training-troubleshooting so the new pages don't repeat material that's already published.

A3b — Training worker process cutover + training UI (TTT)

Goal: Move training out of the backend process into a dedicated worker, so a long AlphaZero Deep run can never starve the API box. Sequenced as 10 PR-sized chunks; each ships independently to dev and each de-risks the next. The three deferred A3a UI items (stacked W/D/L viz, Master-vs-bot split, “no knobs” disclosure) ride the same sprint because they consume data the worker now produces.

PR-sized chunks:

- ☒ **A3b.1 — Spike: new Fly app + BullMQ + hello-world job.** Stands up the queue plumbing without touching any real training behavior; proves Redis reachability + queue lib choice + a separate worker process. Files: backend/src/queue/trainingQueue.js (producer-side Queue + enqueuePing), backend/src/queue/worker.js (BullMQ Worker entrypoint, dispatch by job name, SIGTERM/SIGINT graceful shutdown), backend/src/queue/jobs/ping.js (pure handler, 3 unit tests). New QA_SECRET-gated POST /api/v1/admin/qa/training-worker/ping endpoint mirrors the training-recovery harness pattern so the round-trip can be smoke-tested against any env without a logged-in admin. New backend/fly.training.toml for the xo-training-staging / xo-training-prod apps — shares the backend image, entrypoint switched via [processes] worker. New docker-compose.yml training-worker service. Sidecar fix: both backend/Dockerfile and backend/Dockerfile.dev now copy package-lock.json and drop --no-package-lock — without the lockfile the workspace install was silently dropping bullmq transitives (tslib / msgpackr / node-abort-controller / cron-parser). Local verification: POST /qa/training-worker/ping → BullMQ → worker → handler → completed at ~7 ms enqueue-to-consume lag.
- ☐ **A3b.2a — Worker can consume training:start end-to-end (shadow mode).** Worker process boots, listens to a training:start job on the queue, and runs the existing _runTraining against a non-prod test skill (the same qa-recovery-seed fixtures the harness already uses). Backend continues to run the loop in-process for real traffic — this PR only proves the worker side. Tests: worker happy path on test fixture; pause/resume + checkpoint write still land at the same per-episode cadence.
- ☐ **A3b.2b — Cut over: backend enqueues instead of setImmediate.** Single flag flip behind a ml.useWorker SystemConfig kill switch for instant rollback. mlService.startTraining switches from setImmediate(() => _runTraining(...)) to queue.add('training:start', ...). Pause / resume / cancel signals route via Redis pub-sub. Crash recovery: BullMQ's job lock + heartbeat replaces resumeOrphanedSessions()'s boot-time scan. SSE training:<sessionId> channel: worker publishes progress events to Redis, backend subscribes and fans out (clients see no contract change). The 3-session parallel crash-recovery harness now exercises the worker path. **Highest-risk PR in A3b** — feature-flag gated so a single SystemConfig flip reverts without a deploy.
- ☐ **A3b.3 — Cap enforcement + rate limits at the queue.** BullMQ limiter: { max: N, duration: 1000 } for global concurrency; per-user groupKey so the existing _enforceUserConcurrencyCap runs at enqueue time, not deep inside mlService. Admin override (ml.maxActiveSessionsForAdmin) preserved through the new path. Tests: cap trip + admin override via queue + frontend variant.
- ☐ **A3b.4 — Graceful shutdown + retry / dead-letter.** Worker SIGTERM handler coordinates with the running training loop: signals pauseSession, awaits checkpoint flush, then closes the BullMQ worker. BullMQ retry policy (transient → auto-retry with backoff; N permanent failures → DL set). New GET /admin/training/dead-letter returns the DL contents + a per-job diagnostic dump (session.summary + last 100 log lines). Tests: rolling restart loses no work (in-flight job

resumes on the next instance); permanent failure lands in DL set.

- ❑ **A3b.5 — Observability.** Per-worker CPU/RAM metrics (`process.resourceUsage()` sampled into the existing resource-counter pipeline). Queue depth gauge + alert ($>N$ waiting for $>5\text{min}$). Dead-letter alert (non-empty). Hook into `Observability_Plan.md`. Bull-Board deferred unless we hit a real need — the existing admin Health page can render the gauge inline.
- ❑ **A3b.6 — Stage soak + promote.** `fly` launch the `xo-training-staging` app (already configured in `fly.training.toml`); wire to staging Redis/DB. Run the 3-session parallel crash-recovery test against the worker path. Load test: multiple concurrent users $\times$ mix of algorithms (`rule_based` + `qllearning` + `dqn` + `alphazero`). Promote to prod via the normal /promote path once green.
- ❑ **A3b.7 — Stacked W/D/L visualization (moved from A3a).** Frontend stacked-area chart per opponent curve, reading from the existing `TrainingMetric` rows the A3a multi-curve eval already writes. Saturated curves (e.g., 100% wins against Easy) auto-fade so attention stays on the curves that are still changing. Tests: chart renders for a session with multiple curves; saturation fade triggers when a curve hits 100% for ≥ 3 consecutive eval points. Can ship in any order after A3b.2b.
- ❑ **A3b.8 — “No knobs” disclosure (moved from A3a).** v1 surface stays presets-only (Quick / Standard / Deep). Advanced controls (custom iterations, learning rate, eval budget overrides) live behind an “Advanced” disclosure that only appears for admins or via a per-user feature flag (`features.trainingAdvancedKnobs`). Tests: presets-only by default; advanced reveals for admin; flagged users see it too.
- ❑ **A3b.9 — Master-vs-bot split stub (moved from A3a; lights up in Phase B).** If eval includes Master, split into “as first-mover” and “as second-mover” sub-charts. No-op for TTT (no Master tier) — wires up for real when Connect 4 ships its Master solver. Tests: shows for sessions with Master metrics, hidden otherwise.

A4 — Multi-skill bot UI + auto-clone

Goal: Finish the Phase 3.8 remainder. By the time Connect 4 ships, multi-game bot management is already in users’ hands.

- ❑ Bot detail page: per-skill cards listing every `BotSkill` row owned by the bot, with per-game ELO and last-trained timestamps.
- ❑ Add-skill flow: from a bot, pick a game the bot doesn’t have a skill for, pick an algorithm, optionally clone hyperparams from an existing skill. Wire to `repointBotPrimarySkill` so `User.botModelId` updates correctly.
- ❑ Auto-clone non-primary skills on rename/rebrand: when the bot is renamed, every skill carries along (already true; verify and test).
- ❑ Picker disambiguation: when a bot has multiple skills, the picker shows “BotName (game)” with the active `gameId`, so users always know what they’re queuing.
- ❑ Profile + Gym nav: deep-link straight to a skill’s training page from the bot card.
- ❑ Pre-C4 validation: with only TTT live, the UI is intentionally lightly populated. We’re validating the data plumbing, not showcasing breadth.
- ❑ Tests: add-skill happy path, rename across multiple skills, picker disambiguation, primary-skill repoint after training completes.

Phase A exit criteria

- ❑ TTT plays identically end-to-end (casual + ranked + tournament) on the new routes and new slug.
- ❑ TTT match-level ELO produces sensible movements for the existing user base.
- ❑ All 5 TTT algorithms train through the worker process with W/D/L streaming, multi-curve eval, and checkpoint resume.
- ❑ Multi-skill bot UI lets a user add, manage, and queue from multiple skills (limited to one game today).
- ❑ Help corpus updated: any TTT-facing docs that reference the old slug, single-game ranked, or the old training UX are refreshed.
- ❑ No open regressions on the V1 acceptance script.

Phase B — Connect 4 implementation

Five sprints. Each builds atop the architecture established in Phase A. If any sprint here surfaces an SDK gap, the work pauses, the SDK is extended (with `Game_SDK_Developer_Guide.md` updated), and the sprint resumes — game logic never leaks into the platform.

B1 — packages/game-connect-four/ package

Goal: A standalone, SDK-conformant game package. Engine + bot interface + tests, no UI yet.

- ☐ Package scaffold: packages/game-connect-four/ with the same shape as packages/game-xo/.
- ☐ Engine: 6×7 board, gravity mechanic, win detection (horizontal/vertical/both diagonals), draw detection (board-full), legal-move enumeration (top-of-column).
- ☐ State serialization: compact format suitable for replay storage and bot weight inputs.
- ☐ meta export: id: 'connect-four', inputMode: 'column', matchFormat: { ranked: 'bo2', tournament: 'bo3', master: 'bo2' }, masterStrategy: 'solver', builtInBots: [easy, medium, hard, master].
- ☐ botInterface:
 - Built-in minimax tiers (Easy depth-2, Medium depth-4, Hard depth-6 to depth-8, Master = the solver).
 - Training hooks for the five algorithms (delegating to the shared @xo-arena/ai engines where game-agnostic; new code only where C4 specifics matter).
- ☐ Master solver: perfect-play minimax with full-depth search and the known solved-game heuristics (center-first opening, threat parity). Off-ladder.
- ☐ Unit tests: engine correctness (win/draw/illegal-move), serialization round-trip, minimax depth correctness, solver vs known opening lines.

B2 — C4 play surface (desktop + mobile)

Goal: Casual quick match against built-in Easy/Medium/Hard works end-to-end via /games/connect-four/...

- ☐ Game component: 6×7 board with Red/Yellow tokens, drop animation, win highlight, draw indicator.
- ☐ Desktop input: click on column to drop.
- ☐ Mobile input: tap-to-preview (semi-transparent token at the top of the column showing where the piece will land), tap-to-confirm (drops the piece). Tap a different column to switch preview. Long-press cancels.
- ☐ Audio cues: drop, win, draw — reusing the existing sdk.playSound() pattern with new sample IDs.
- ☐ Theming: Red/Yellow color scheme conforms to platform tokens (packages/sdk theme exports).
- ☐ Spectate: rendered-table mode works; piece drops animate for spectators.
- ☐ Replay: replays of C4 games render via the same replay infrastructure as TTT.
- ☐ Bot vs human: server-side bot moves work for built-in Easy/Medium/Hard. Master is not yet wired here (Phase B5).
- ☐ Tests: e2e casual match vs each built-in tier on desktop + mobile viewports.

B3 — C4 ranked + tournament

Goal: Ranked best-of-2 and tournament best-of-3 work for C4 with zero new infrastructure — A2's match layer just lights up.

- ☐ Wire C4 into ranked match orchestration. Alternating colors, ELO updates per match.
- ☐ Wire C4 into the tournament service. Best-of-3 with random-color game 3 tiebreaker.
- ☐ Classification ladder: C4 ELO tracked separately from TTT (already supported by GameElo (userId, gameId) unique constraint).
- ☐ Landing nav: C4 appears in the games picker; BotFilterBar shows both games.
- ☐ Tests: ranked best-of-2 happy path, tournament with 1-1 game-3 tiebreaker, classification ladder shows independent TTT and C4 ratings for the same user.

B4 — C4 training, all five algorithms

Goal: Every algorithm trains through the A3 architecture with no platform-side changes.

- ☐ Rule-Based: heuristic rules (center preference, win-on-move, block-on-move, fork detection, edge play). Zero-time “training.”
- ☐ Minimax (Quick Bots): tier-label bump only, same pattern as TTT Quick Bots — user:<id>:minimax:<tier> for C4.
- ☐ MCTS: classical UCB1 with rollouts. Tunes simulation count, exploration constant, optional rollout policy.
- ☐ DQN: experience replay + target network. Standard preset hyperparams.
- ☐ AlphaZero: policy/value network with PUCT-guided MCTS. Three presets:
 - Quick — ~30 minutes, small network, low simulation count. Auto-confirmed.
 - Standard — ~4 hours, medium network. **Hits admin-approval threshold.**
 - Deep — ~40 hours, full network + high sim count. Always admin-approved.

- ☐ Eval curves: primary opponent per algorithm (e.g., AZ evals against Hard minimax). Easy + Medium side curves render saturated and auto-fade.
- ☐ Streaming UX: TTT users seeing W/D/L curves in realtime today will see the same UI for C4 with no relearning.
- ☐ Tests: each algorithm reaches expected milestones from a known seed in CI-bounded episode counts.

B5 — Master tier + Master Challenge

Goal: Perfect-play Master, off-ladder, with the Master Challenge user flow.

- ☐ Master button on the C4 home: opens a best-of-2 Master Challenge (one game as first-mover, one as Master first-mover).
- ☐ Master plays the perfect-play solver from B1. Move latency budgeted (solver is fast enough at depth-required for solved-game optimality, but cache opening lines).
- ☐ No ELO update on Master matches. Per-bot “vs Master” stat block on the bot detail page (wins / draws / losses).
- ☐ Training integration: AlphaZero bots can opt into “vs Master” eval as an additional curve (showing first-mover-draw rate over time — the AlphaZero teaching moment).
- ☐ Landing UI: Master section on C4 home page explains the off-ladder rule and what “beating Master” means (linking to `Help_Corpus/solved-games-and-master.md`).
- ☐ Tests: Master Challenge happy path, off-ladder ELO behavior (no rating movement), AZ training with Master-vs-bot eval curve.

Phase B exit criteria

- ☐ C4 ships through `/games/connect-four/...` for casual, ranked, tournament, and Master Challenge.
 - ☐ All five algorithms train through the A3/A3b architecture. AZ Deep on C4 succeeds at least once on staging without affecting API latency.
 - ☐ Multi-skill bots can hold both a TTT and a C4 skill, each with independent ELO.
 - ☐ Replay, spectate, ranked classification all work for C4 at parity with TTT.
 - ☐ Corpus matches reality — the Connect 4 pages in `Help_Corpus/` reflect what shipped.
-

Phase C — Polish

Approximately 3-4 sprints. Lower urgency than Phase B but completes the end-to-end experience.

C1 — Journey + coaching updates

- ☐ Update the onboarding journey to surface a Connect 4 milestone after the user’s first ranked TTT win.
- ☐ Add coaching cards for C4 specifics (column thinking, first-mover advantage, threat parity).
- ☐ Update `Intelligent_Guide_Requirements.md` and `Intelligent_Guide_Implementation_Plan.md`.

C2 — Training UX completions

- ☐ Stop-early button on active sessions (with confirmation + “save current checkpoint”).
- ☐ 7-day rollback button on bot detail page: revert a skill’s weights to a previous checkpoint within a 7-day retention window.
- ☐ Advanced disclosure (v1.1): expose hyperparameter knobs behind an “Advanced” toggle for users who want to tune past presets.

C3 — Corpus completions

- ☐ Write the queued cost/preset docs: “What each preset does”, “How long does training take?”, “What does training cost?”, “What happens if I close the tab?”, “Stop training early”, “Why no knob tweaking”, “Why AZ slow on C4 fast on TTT”.
- ☐ Refresh `bots-overview.md`, `gym-train-tab.md`, `understanding-training-charts.md` with C4-specific examples.
- ☐ Render PDFs via the tuned pandoc invocation.

C4 — Accessibility + mobile polish

- ☐ Keyboard navigation for the C4 board (arrow keys + enter to drop).
- ☐ Screen-reader labels for board state (“Column 4, row 3, your turn”).

- Mobile UX refinement pass — confirm the tap-to-preview-tap-to-confirm gesture works cleanly across viewports and doesn't conflict with browser swipe gestures.

C5 — Observability sweep

- C4-specific dashboards: completion rate per tier, AZ Master-draw-rate over time, training queue depth, worker CPU/RAM under typical load.
- Alert thresholds: dead-letter queue depth, training session failures per hour, API latency regression during heavy training.

C6 — Legacy slug cleanup

Goal: Drop the `xo` → `tic-tac-toe` alias once Phase B has shipped and external callers have had multiple release cycles to migrate. Cleanup, not architecture — sequenced here on purpose so the deprecation window is long enough that nothing in the wild still calls the old slug.

Pre-conditions: Phase B is shipped to production. No staging or production access logs show requests with `gameId=xo` or path `/play/xo*` for at least one release cycle (≥ 2 weeks). Verify via the request logs / observability dashboards before starting.

- Add 301 Moved Permanently redirects in `landing/server.js` and (if needed) the backend for legacy paths: `/play/xo*` → `/games/tic-tac-toe/play*`, `/xo/*` → `/games/tic-tac-toe/*`. Keep the redirects in place for one further release.
- Remove `LEGACY_SLUG_MAP` from `backend/src/constants/games.js` and `tournament/src/constants/games.js`. `resolveGameSlug('xo')` should now return null (caller's ?? `rawGameId` fallback then surfaces “Unknown game slug” 404 via `validateGameSlug`).
- Sweep for any residual 'xo' literals in code (excluding migration files, historical perf-trace doc citations, and `Help_Corpus` search-alias tags). Update or delete each one.
- Update `Game_SDK_Developer_Guide.md` to drop the “legacy alias accepted” caveat — kebab-case canonical slugs only.

Risks + mitigations

Risk	Mitigation
Slug-rename migration corrupts production data. Renaming <code>xo</code> → <code>tic-tac-toe</code> across <code>BotSkill</code> , <code>GameElo</code> , <code>Tournament</code> is a write-heavy migration.	Dry-run on staging with a prod data snapshot. Migration runs inside a transaction with explicit row counts logged. Rollback plan: a reverse-rename migration is staged before deploy.
Match-level ELO produces unexpected user-visible drops or jumps at cutover. Even with frozen historical rows, the first few matches under the new formula will move ratings under different math.	Communicate proactively in release notes. The new math gives smaller per-match deltas (it's “match score 0.5” vs “game-loss 0.0”), so the practical effect is gentler movement. Monitor for outliers.
AlphaZero Deep run on C4 starves the API box despite the worker. Per-user cap and worker isolation should prevent this, but a misconfigured Fly app could fall back to in-process.	Pre-launch load test on staging: queue a real AZ Deep run while a separate user plays ranked. API p95 latency must not regress.
SDK gap surfaces mid-Phase B. Discovering during B2 that the SDK can't express column-input semantics would block the sprint.	Phase A1 audits explicitly for this. If a gap appears in B, the SDK is extended in place; we don't ship workarounds.
Master solver is too slow at depth-required. Perfect-play C4 needs deep search; if move latency is too high, the Master tier feels broken.	Cache the entire opening book (first 10-12 plies) as a lookup table. Worst case, fall back to a precomputed move table for the opening, switch to live search after the book exits.
Mobile column-input gesture conflicts with browser scroll. A column tap shouldn't scroll the page.	touch-action: manipulation on the board; e2e tests on mobile viewports.
Worker deployment introduces a new failure surface. A worker outage could halt all training.	Status page reflects worker health. If the worker is down, new training-start requests return a clear “training temporarily unavailable” error and queue UI surfaces the outage. Existing ranked/casual/tournament play is unaffected (those don't go through the worker).

Sequencing summary

Phase A (TTT only):

- A1 – SDK + routing + slug rename
- A2 – Match-based play + match-level ELO
- A3a – Training data + UX rework (in-process)
- A3b – Worker process cutover
- A4 – Multi-skill bot UI + auto-clone

Phase B (Connect 4 build):

- B1 – packages/game-connect-four/
- B2 – C4 play surface (desktop + mobile)
- B3 – C4 ranked + tournament
- B4 – C4 training, all five algorithms
- B5 – Master tier + Master Challenge

Phase C (polish):

- C1 – Journey + coaching
- C2 – Training UX completions
- C3 – Corpus completions
- C4 – Accessibility + mobile polish
- C5 – Observability sweep

Phases ship sequentially. Sprints within a phase can run in parallel where they don't have ordering dependencies (e.g., B2 and B3 can overlap; B4 needs the C4 package from B1; B5 needs the C4 surface from B2).

Open items (will be answered as the plan executes)

- **Worker pool sizing.** Start with count=1 on staging, count=2 on prod, scale based on observed queue depth. Final sizing TBD post-launch.
- **AZ C4 Deep economics.** Per the corpus, AZ C4 Deep is roughly \$3–50 per run depending on Fly machine class. The exact admin-approval UX (credits charged? burned against an admin pool?) is open — pin down before Phase B4.
- **Master solver move-cache size.** Opening-book lookup table size depends on how deep the book runs. Spike during B1.
- **Color naming.** Red/Yellow is canonical; settle whether colorblind-mode swaps to a high-contrast pair or uses pattern fills. Decide during B2.