AI Arena — Bot Training Guide

Getting Started: Bots, skills & the Gym

Every bot in AI Arena learns through one or more **skills**. A skill is an AI model tied to a specific game — its learning algorithm plus the trained weights. As AI Arena expands across multiple games (XO today, Pong and others to come), a bot can carry **multiple skills** — one per game — each trained independently.

Before you can train, you need to **create a bot** on your Profile page, then **add a skill** for the game you want to play. Adding a skill means choosing an algorithm — the method by which the skill learns from experience.

AI Arena offers six algorithms, ranging from classic tabular methods to neural-network approaches:

Algorithm	Style	Best For
Q-Learning	Tabular, off-policy	Fast, reliable, great starting point
SARSA	Tabular, on-policy	Conservative, less exploitable play
Monte Carlo	Tabular, episode-level	Clean credit assignment, no bias
Policy Gradient	Tabular softmax	Unpredictable, human-like play
DQN	Neural network	Generalises unseen board states
AlphaZero	Neural network + MCTS	Strongest ceiling, lookahead search

Once your bot has a skill, you bring it to the **Gym** to train. Training runs in sessions — you pick a mode (self-play, vs minimax, alternating), configure the exploration settings, and run a batch of episodes. Each episode is a full game your bot plays and learns from. Over many sessions you build up a trained, competitive skill ready to challenge others on the leaderboard.

This guide is your training manual. It explains the core concepts that apply to all algorithms, then goes deep on each one — what it does differently, which settings to use, and session-by-session recipes for getting the best results. The Benchmarking and Troubleshooting sections at the end help you measure progress and diagnose problems.

Table of Contents

Core Concepts

Epsilon (ε) — Exploration vs. Exploitation

Epsilon controls the probability of choosing a **random move** instead of the agent's current best move.

- **ε = 1.0** — fully random (pure exploration, learns nothing yet)
- **ε = 0.5** — 50/50 mix
- **ε = 0.05** — mostly greedy, small residual exploration
- **ε = 0.0** — fully greedy (exploitation only, no learning)

Epsilon starts high and decays over time. Decaying too fast locks in a bad policy before it has converged. Decaying too slow wastes episodes on random play once the agent already knows the right moves.

Epsilon Decay Methods

Method	Formula	Best For
Exponential	`ε *= decay_rate` each episode	Most algorithms (default)
Linear	ε decreases evenly over N episodes	DQN, fixed budget
Cosine	Cosine curve from start → min	Smoother warmup/cooldown

Rule of thumb: ε should reach `epsilonMin` at roughly **80% of total planned episodes**, leaving the last 20% for pure exploitation to lock in the learned policy.

Epsilon Decay Rate Quick Reference

> *In the app, this is the "**Rate**" field (visible when "Decay schedule" = Exponential).*

Rate	Episodes to reach ε=0.05 (starting at ε=1.0)
0.995	~590
0.999	~2,990
0.9995	~5,990
0.9999	~29,950
0.99995	~59,900

Training Modes

Mode	Description	Best For
Self-play	Bot plays both sides simultaneously	All algorithms, best for tabular methods
vs Minimax	Bot plays against a deterministic opponent	Curriculum training, final polish
Alternating	Bot plays X one episode, O the next	Ensures symmetry for both sides

Curriculum Training

Available in vs-Minimax mode. The system automatically advances difficulty (novice → intermediate → advanced → master) once the bot wins 65% of the last 100 games. Start at **novice** — jumping straight to advanced wastes early episodes against an opponent the agent cannot yet learn from.

Gamma (γ) — Discount Factor

Controls how much future rewards are valued vs immediate ones.

- **γ = 0.9** — standard, good for tic-tac-toe (short games)
- **γ = 0.99** — near-sighted future; used for AlphaZero
- Lower γ → agent plays for quick wins; higher γ → agent plans further ahead

For tic-tac-toe specifically, γ = 0.9 is ideal since games are at most 9 moves.

> ***In the app:*** *γ is exposed in the* ***DQN Train tab*** *as the* ***"Gamma (γ)"*** *dropdown (0.85 / 0.90 / 0.95 / 0.99). For all other algorithms it is hardcoded (0.9 for tabular, 0.99 for AlphaZero). For tabular methods only, γ can also be swept via the* ***Auto-Tuner*** *tab.*

Learning Rate (α)

Controls how aggressively Q-values are updated per step.

- **Tabular methods (Q-Learning, SARSA, MC):** α = 0.3–0.5 is safe. Higher is faster but noisier.
- **Neural net methods (DQN, AlphaZero):** α = 0.001. Too high causes divergence.

> ***In the app:*** *α is* ***not exposed in the Train tab*** *for any algorithm. It is hardcoded at 0.3 (tabular) and 0.001 (neural net). For tabular methods only, α can be swept via the* ***Auto-Tuner*** *tab.*

↑ Back to top

Algorithm Overview

Algorithm	Type	State Space	Episodes to Competent	Best Strength
Q-Learning	Tabular, off-policy TD	Explicit table	3,000–8,000	Fast, reliable, debuggable
SARSA	Tabular, on-policy TD	Explicit table	4,000–10,000	Conservative, safe policy
Monte Carlo	Tabular, episode-level	Explicit table	5,000–15,000	No bootstrapping bias
Policy Gradient	Tabular softmax	Explicit table	8,000–20,000	Stochastic, hard to exploit
DQN	Neural net, off-policy	Neural network	10,000–25,000	Generalizes unseen states
AlphaZero	Neural net + MCTS	Neural network	5,000–15,000	Strongest ceiling

If you want the strongest bot: AlphaZero > DQN > Q-Learning. **If you want the fastest to train:** Q-Learning > SARSA > Monte Carlo.

↑ Back to top

Q-Learning

Algorithm: Off-policy TD control. Updates Q(s,a) immediately after each move using the best possible next action (regardless of what will actually be played).

State space: Explicit table keyed by board string. Tic-tac-toe has ~5,478 reachable states — Q-Learning covers this completely within a few thousand episodes.

Recommended Settings

UI Label	Value	Configurable?	Notes
Decay schedule	**Exponential**	Train tab	
Rate	**0.995**	Train tab (Exponential only)	Reaches ε≈0.05 by ~590 episodes
Epsilon min	**0.05**	Train tab	5% residual exploration
Reset ε to 1.0 at start	**checked**	Train tab	Always start fully random
Learning rate (α)	0.3	**Hardcoded**	Not shown in UI — Auto-Tuner can sweep it
Discount factor (γ)	0.9	**Hardcoded**	Not shown in UI — Auto-Tuner can sweep it

Session Recipe

Session	Episodes	Mode	Opponent	Purpose
1	2,000	Self-play	—	Bootstrap all states
2	3,000	Self-play	—	Converge Q-table
3	2,000	vs Minimax	Curriculum (novice→master)	Polish vs deterministic play

Total: ~7,000 episodes

Expected Results by Episode Count

Episodes	vs Random	vs Easy Minimax	vs Hard Minimax
500	50–60%	20–30%	<5%
2,000	75–85%	50–65%	10–20%
5,000	85–92%	70–80%	30–50%

8,000+	90–95%	80–90%	50–65%

Tips

- Q-Learning is **off-policy**: it learns the greedy strategy even while exploring. This makes it faster than SARSA at reaching a strong policy.
- Self-play is ideal for session 1 because it trains both X and O simultaneously.
- Once the bot reaches ~80% vs random, switch to curriculum minimax to sharpen strategy.
- If the Q-table stalls (win rate plateaus for 2,000+ episodes), lower α to 0.1 and run 2,000 more episodes — aggressive learning rates overwrite good values at low epsilon.

↑ Back to top

SARSA

Algorithm: On-policy TD control. Updates Q(s,a) using the action that will *actually* be taken next (not the greedy max). This makes SARSA more conservative — it learns a safer policy that accounts for its own exploration behavior.

State space: Same explicit table as Q-Learning.

Recommended Settings

UI Label	Value	Configurable?	Notes
Decay schedule	**Exponential**	Train tab	
Rate	**0.995**	Train tab (Exponential only)	
Epsilon min	**0.05**	Train tab	Keep slightly higher than Q-Learning
Reset ε to 1.0 at start	**checked**	Train tab	
Learning rate (α)	0.3	**Hardcoded**	Not shown in UI — Auto-Tuner can sweep it
Discount factor (γ)	0.9	**Hardcoded**	Not shown in UI — Auto-Tuner can sweep it

Session Recipe

Session	Episodes	Mode	Opponent	Purpose
1	3,000	Self-play	—	Explore all states under SARSA's cautious policy
2	3,000	Self-play	—	Converge
3	2,000	vs Minimax	Curriculum	Polish

Total: ~8,000 episodes

Expected Results

Episodes	vs Random	vs Easy	vs Hard
1,000	45–55%	15–25%	<5%
4,000	70–80%	55–65%	15–25%
8,000+	85–90%	70–80%	40–55%

SARSA vs Q-Learning

- SARSA converges **~20% slower** but produces a policy that is less exploitable in practice.
- The gap between the two narrows as epsilon decays to 0 — at inference time (ε=0) both become identical greedy agents.
- Prefer SARSA if you want a bot that draws rather than gambles; prefer Q-Learning if you want raw win rate.

↑ Back to top

Monte Carlo

Algorithm: Every-visit MC control. Waits until the end of the episode, then propagates the actual discounted return backward through all visited (state, action) pairs. No bootstrapping — every update is based on real experienced outcomes.

Convergence: Slower than TD methods because updates only happen at episode end, not after each move. However, updates are unbiased — there is no bootstrapping error accumulating over time.

Recommended Settings

UI Label	Value	Configurable?	Notes
Decay schedule	**Linear**	Train tab	Linear works well for MC; more even coverage
Epsilon min	**0.05**	Train tab	
Reset ε to 1.0 at start	**checked**	Train tab	
Learning rate (α)	0.2	**Hardcoded**	Not shown in UI — lower than TD because MC updates are noisier; Auto-Tuner can sweep it
Discount factor (γ)	0.9	**Hardcoded**	Not shown in UI — Auto-Tuner can sweep it

> ***Note:*** *Monte Carlo does not show a "Rate" field when Linear decay is selected — the decay spreads evenly across the session's episode count automatically.*

Session Recipe

Session	Episodes	Mode	Opponent	Purpose
1	5,000	Self-play	—	Wide exploration — many unique trajectories
2	5,000	Self-play	—	Reinforce good trajectories
3	3,000	vs Minimax	Curriculum	Tighten strategy

Total: ~13,000 episodes

Expected Results

Episodes	vs Random	vs Easy	vs Hard
2,000	40–55%	15–25%	<5%
6,000	65–75%	45–60%	10–20%
12,000+	80–88%	65–75%	30–45%

Tips

- Use **linear decay** — Monte Carlo benefits from staying exploratory longer because each episode covers a unique trajectory. Exponential decay collapses exploration too early.
- MC is the most data-hungry algorithm but has the cleanest credit assignment: every move in a winning game is credited, every move in a losing game is penalized.
- The final policy quality is often comparable to Q-Learning at 1.5–2x the episode count.

↑ Back to top

Policy Gradient

Algorithm: REINFORCE with a tabular softmax policy. Instead of learning Q-values and acting greedy, it directly learns action preferences θ(s,a) and samples from a probability distribution. This produces naturally stochastic play — the bot doesn't always make the same move in the same position.

Key difference from Q-Learning: The policy never fully "locks in" a deterministic move. At inference, it picks the highest-preference action, but the learned distribution means it can be less predictable against human opponents.

Recommended Settings

UI Label	Value	Configurable?	Notes
Decay schedule	**Cosine**	Train tab	Smooth schedule suits PG's gradient updates
Epsilon min	**0.05**	Train tab	
Reset ε to 1.0 at start	**checked**	Train tab	Controls fallback to random; PG uses softmax sampling internally

Learning rate (α)	0.01	**Hardcoded**	Not shown in UI — PG is sensitive; hardcoded lower than tabular default; Auto-Tuner can sweep it
Discount factor (γ)	0.9	**Hardcoded**	Not shown in UI — Auto-Tuner can sweep it

> ***Note:*** *Policy Gradient does not show a "Rate" field when Cosine decay is selected.*

Session Recipe

Session	Episodes	Mode	Opponent	Purpose
1	5,000	Self-play	—	Build preference distribution across states
2	5,000	Self-play	—	Reinforce high-return trajectories
3	5,000	vs Minimax	Curriculum	Sharpen vs deterministic opponent
4	3,000	vs Minimax	Master	Final polish

Total: ~18,000 episodes

Expected Results

Episodes	vs Random	vs Easy	vs Hard
2,000	35–50%	10–20%	<5%
8,000	60–72%	40–55%	10–20%
15,000+	75–85%	60–72%	25–40%

Tips

- Policy Gradient is the **slowest tabular algorithm** but produces the most unpredictable play.
- The learning rate is critical: α > 0.05 causes oscillation. If the win rate bounces wildly between sessions, halve α.
- The cosine decay schedule works particularly well here — it mirrors the natural "explore then refine" lifecycle of REINFORCE.
- PG bots make better opponents for human players because they are not fully deterministic.

↑ Back to top

DQN (Deep Q-Network)

Algorithm: Neural network function approximation for Q-values. Instead of an explicit state table, a small MLP `[9 → hidden → 9]` represents Q-values. Uses experience replay (circular buffer) and a separate target network for stability.

Key advantages:

- Can generalize across similar board states that were never exactly seen during training.
- No state table — model size stays fixed regardless of states visited.
- With Adam optimizer and the corrected adversarial Bellman equation, converges reliably.

DQN Train tab controls:

UI Label	Default	Description	Configurable?
Batch	32	Samples per gradient step	Yes — Train tab
Replay buffer	10,000	Experiences stored for random sampling	Yes — Train tab
Target update	100	Steps between target network syncs	Yes — Train tab
Network architecture	[32]	1–3 hidden layers, each 8/16/32/64/128 neurons	Yes — Train tab
Gamma (γ)	0.90	Discount factor (0.85 / 0.90 / 0.95 / 0.99)	Yes — Train tab
Decay schedule	Exponential	Epsilon decay curve	Yes — Train tab
Rate	0.9999	Per-episode decay multiplier (Exponential only)	Yes — Train tab
Epsilon min	0.05	Minimum epsilon floor	Yes — Train tab
Reset ε to 1.0 at start	checked	Whether to restart exploration	Yes — Train tab
Learning rate (α)	0.001	Adam optimizer step size	**Hardcoded** — not exposed

> ***Architecture changes reset weights.*** *If you change the layer layout (e.g., [32] → [64, 64]) the Train tab will show an amber warning and the existing trained weights will be discarded — training starts fresh with the new shape. The model's stored architecture is updated to match after the session completes.*

Recommended Settings

UI Label	Value	Notes
Network architecture	**[64, 64]**	Two layers; best quality/speed tradeoff for tic-tac-toe
Gamma (γ)	**0.95**	Better than default 0.90 — plans further ahead in 9-move games
Decay schedule	**Exponential**	Required for multi-session training — see note below
Rate	**0.9999**	Reaches ε≈0.05 by ~29,950 episodes across all sessions
Epsilon min	**0.05**	

Reset ε to 1.0 at start	**checked**	
Batch	**32**	Default; increase to 64 for [128, 128] nets
Replay buffer	**10,000**	Good for 30k-episode runs
Target update	**100**	Default; do not lower below 50

> ***Why Exponential for DQN?*** *Linear and Cosine decay spread ε from its current value to Epsilon min evenly across one session's episodes — the Rate field is hidden because it isn't used. For a 10k-episode session starting at ε=1.0 with Epsilon min=0.05, Linear would reach the floor at episode 10,000, leaving sessions 2–4 with no exploration. Exponential at Rate 0.9999 carries exploration smoothly across 30,000+ episodes, matching the ε estimates in the session recipe below. Use Linear/Cosine only for a single all-in-one session.*

Session Recipe

The **ε start → end** column shows predicted values based on Exponential 0.9999 decay — these are not inputs you configure. ε start is determined by the previous session's stored end state. The only epsilon control in the Train tab is the **"Reset ε to 1.0"** checkbox:

- **Session 1** — check it (start fresh at 1.0)
- **Sessions 2–4** — uncheck it (continue from the stored value left by the previous session)

ε end is likewise not settable; it is simply where epsilon lands after N episodes of decay at your chosen Rate.

Session	Episodes	Mode	Opponent	ε start → end (predicted)	Reset ε to 1.0	Purp
1	10,000	Self-play	—	1.0 → ~0.37	**checked**	Warm replay buffer learn l move validit
2	10,000	Self-play	—	~0.37 → ~0.14	unchecked	Core policy learni
3	10,000	vs Minimax	Curriculum (novice→advanced)	~0.14 → ~0.05	unchecked	Strate refine
4	5,000	vs Minimax	Master	~0.05 → ~0.05	unchecked	Final harde

Total: ~35,000 episodes

Expected Results

Episodes	vs Random	vs Easy	vs Hard

5,000	50–65%	20–35%	<10%
15,000	72–82%	50–65%	20–35%
25,000	82–90%	65–78%	35–50%
35,000+	88–93%	75–85%	50–65%

Understanding the Replay Buffer

The replay buffer is a **circular memory** that stores past game experiences as `(board_state, action, reward, next_board_state, done)` tuples. During each training step the DQN samples a random mini-batch from this buffer — it does *not* train on the episode it just played.

Why this matters:

- **Breaks temporal correlation** — consecutive episodes are highly correlated (same opening, similar board patterns). Random mini-batch sampling produces more diverse batches, which stabilises gradient descent and prevents catastrophic forgetting.
- **Reuses rare experiences** — important board positions (e.g. fork threats, game-ending states) can be revisited many times even if they are infrequently encountered. Without the buffer each experience is used exactly once and then discarded.

Buffer size relative to episode count:

Episodes trained	Buffer = 10,000	Buffer = 20,000
5,000	Buffer is half full; all experiences available	Buffer is quarter full
10,000	Buffer exactly full; oldest evicted as new arrive	Still holds all 10k experiences
30,000	Only the most recent 10k episodes are kept	Keeps the most recent 20k

A buffer that is **too small** for long runs means the network only learns from recent play — it can forget earlier lessons and oscillate. Increase to 20,000–50,000 for sessions over 20,000 episodes.

A buffer that is **too large** simply has empty slots early in training; this is harmless — the network trains on what is available.

When to tune the replay buffer:

- Win rate oscillates wildly → buffer too small (increase it)
- Bot forgets good moves it knew earlier → buffer too small relative to total episodes trained
- Training speed feels slow on a powerful machine → batch size is the knob to increase, not buffer size

Understanding the Training Charts

The live training panel shows two sets of lines for win/draw/loss rates:

- **Thick solid lines** — the *recent* rate for the current batch of episodes only (resets each update interval)
- **Thin dashed lines** — the *cumulative* average since the session started

The recent lines are the most informative signal for whether the model is actively learning. If a bot is genuinely improving — for example, drawing more in the last 100 episodes than it did in the first 1,000 — the recent draw rate line will climb while the cumulative line rises slowly (weighted down by early losses).

If both lines are flat, the model is not converging in this session.

Network Architecture Guide

Network architecture is now configurable **in the Train tab** using the layer builder. You can add 1–3 hidden layers, each with 8, 16, 32, 64, or 128 neurons. Changing the architecture resets the model's weights — the amber warning in the UI will tell you when this will happen.

Architecture	Parameters	Training Speed	Recommended For
`[32]`	~380	Fastest	Quick experiments
`[64]`	~700	Fast	Good single-layer baseline
`[64, 64]`	~1,350	Medium	**Best quality/speed tradeoff**
`[128, 64]`	~2,700	Slower	Diminishing returns for tic-tac-toe

For tic-tac-toe, `[64, 64]` is the sweet spot — larger networks don't raise the ceiling but take longer to train.

Tips

- Always use **self-play for sessions 1–2**. The replay buffer fills faster and both X and O perspectives are covered simultaneously.
- Use **Exponential decay with Rate 0.9999**. Linear/Cosine decay from ε=1.0 to 0.05 within a single 10k-episode session — leaving nothing for future sessions. Exponential at 0.9999 spreads exploration gradually across all sessions.
- The minimum viable run is **15,000 episodes** to fill and recycle the replay buffer enough for the Bellman targets to stabilize.
- If the win rate is flat after 20,000 episodes: change the architecture to `[64, 64]` in the Train tab (weights reset, but the new capacity helps more than extra episodes on a small net).
- Use **Gamma = 0.95** — it lets the agent plan further ahead without introducing instability in 9-move games.

↑ Back to top

AlphaZero

Algorithm: Monte Carlo Tree Search (MCTS) guided by two neural networks — a policy net `[9→64→32→9]` (softmax output) and a value net `[9→64→32→1]` (tanh output). Each episode runs `numSimulations` MCTS rollouts to build a visit-count distribution, then trains both networks on the collected examples.

Key advantages:

- MCTS guarantees the agent considers multiple lookahead paths before each move — not just immediate reward.
- The value network provides a learned evaluation function.

- No epsilon — exploration is naturally built into the PUCT tree search.
- Consistently produces the strongest policy of all algorithms given sufficient episodes.

AlphaZero Train tab controls:

UI Label	Default	Description	Configurable?
Simulations	50	MCTS rollouts per move. Higher = stronger but slower.	Yes — Train tab
PUCT	1.5	Exploration constant in tree search. Higher = more exploration.	Yes — Train tab
Temperature	1.0	Randomness in move selection from visit counts.	Yes — Train tab
Learning rate (α)	0.001	Shared learning rate for policy + value nets	**Hardcoded**
Discount factor (γ)	0.99	Future reward weighting	**Hardcoded**

> *AlphaZero has **no epsilon** — exploration is naturally built into the PUCT tree search. There is no "Exploration" section for AlphaZero in the Train tab.*

Recommended Settings

UI Label	Value	Configurable?	Notes
Simulations	**100**	Train tab	Double the default for much stronger search
PUCT	**1.5**	Train tab	Default is good; increase to 2.0 for more exploration early
Temperature	**1.0**	Train tab	Reduce to 0.5 after first 5,000 episodes for more decisive play
Learning rate (α)	0.001	**Hardcoded**	Adam optimizer; not exposed in any tab
Discount factor (γ)	0.99	**Hardcoded**	Not exposed; higher than tabular default because AZ plans further ahead
Epsilon / exploration	—	**Not applicable**	AlphaZero has no epsilon — PUCT handles all exploration in the tree

Session Recipe

Session	Episodes	Simulations	Temperature	Purpose
1	3,000	50	1.0	Fast bootstrap — build initial policy/value estimates
2	5,000	100	1.0	Deeper search, refine both nets
3	5,000	100	0.5	More decisive play, sharpen value net

4	2,000	200	0.1	Near-deterministic fine-tuning

Total: ~15,000 episodes

> ***Note:*** *AlphaZero episodes are significantly slower than tabular episodes because each move runs* `numSimulations` *MCTS rollouts. Expect 5–10x more wall-clock time per episode vs Q-Learning.*

Expected Results

Episodes	vs Random	vs Easy	vs Hard
1,000	60–72%	35–50%	10–20%
5,000	78–87%	60–75%	35–50%
10,000	87–93%	75–85%	55–70%
15,000+	92–97%	82–90%	65–80%

AlphaZero reaches competency **much faster per-episode** than other algorithms because MCTS provides strong implicit lookahead even early in training. The episode count is lower but wall-clock time is higher.

Tips

- **Self-play only** — AlphaZero is designed exclusively for self-play. It has no concept of an external opponent function.
- Start with **Simulations=50** to fill the experience buffer quickly, then increase to 100+ as the networks gain accuracy.
- **Temperature** controls how the final move is sampled from MCTS visit counts. At 1.0, proportional sampling adds diversity. At 0.1, the most-visited child is almost always chosen. Reduce Temperature over training for a sharper policy.
- **PUCT=1.5** balances exploitation vs exploration in the tree. If the bot gets stuck in repetitive patterns early, try PUCT=2.0.
- AlphaZero's policy net outputs are probabilities, not Q-values. The Explainability tab shows these as move probabilities — a strong AZ bot will show clear high-probability cells for center and corner openings.

↑ Back to top

Benchmarking & Evaluation

Always run a benchmark after training to get objective scores before deploying a bot.

What the Benchmark Measures

Scenario	Opponent	Win Rate Goal (good bot)
vs Random	Pure random	> 85%
vs Easy	Novice minimax	> 70%
vs Medium	Intermediate minimax	> 55%

vs Tough	Advanced minimax	> 35%
vs Hard	Perfect minimax	> 15% (draws are good here)

> *A perfect tic-tac-toe player can never lose — only win or draw. "vs Hard" scores above 15% indicate the bot is finding genuine wins against a perfect opponent by exploiting first-move advantage.*

Benchmark-to-ELO Correlation

Avg Benchmark Win Rate	Approximate ELO
< 30%	< 900
30–45%	900–1,050
45–60%	1,050–1,200
60–72%	1,200–1,400
72–82%	1,400–1,600
> 82%	1,600+

Head-to-Head: When to Use

Use head-to-head (H2H) to compare two candidate models trained with different configs. Run at least **200 games** for statistical significance. The ELO update from H2H reflects the true relative strength better than benchmark scores against fixed opponents.

p-Value Interpretation

Every benchmark result includes a p-value testing whether win rate > 50%.

- **p < 0.05** — the result is statistically significant (bot is genuinely performing above chance)
- **p > 0.10** — sample size too small or the bot is near 50% — run more games or more training

↑ Back to top

Troubleshooting

Win rate is flat / not improving

1. **Rate (epsilon decay) too fast** — check what epsilon is at the stall point. If ε < 0.1 and the bot hasn't converged, restart with a slower Rate (e.g. 0.999 → 0.9999).
2. **Learning rate too high** — for tabular methods: use the Auto-Tuner to find a better α. For DQN: α is hardcoded at 0.001 (well-tuned, not the issue).
3. **Opponent mismatch** — don't train vs Hard minimax until the bot reliably beats Easy. Use curriculum.
4. **Too few episodes before vs-minimax** — always do self-play first to bootstrap all states/weights before switching to a fixed opponent.

Win rate is oscillating wildly

1. **Learning rate too high** — for tabular methods: use the Auto-Tuner to sweep α. For DQN: α is fixed at 0.001 (this isn't the cause).
2. **DQN only: Replay buffer too small** — if "Replay buffer" < 5,000 and you're doing 10k+ episodes, old experiences are evicted before they can be replayed enough. Increase to 20,000.
3. **DQN only: Target update too low** — syncing the target network every 10 steps instead of 100 causes instability. Keep "Target update" at 100+.

Q-Learning / SARSA / MC win rate hits 80% then never improves

The tabular Q-table is fully converged. Switching algorithms (to DQN or AlphaZero) is the right move — not more episodes of the same algorithm.

AlphaZero is very slow

Each episode runs "Simulations" MCTS passes, each a full tree traversal. At Simulations=200 this is ~200× slower than a Q-Learning episode. Start with 50, confirm the bot is learning (rising win rate on the analytics tab), then increase.

Bot plays well as X but poorly as O (or vice versa)

The bot was trained asymmetrically. Run 2,000–3,000 additional self-play episodes using **alternating** mode so both X and O perspectives receive balanced updates.

DQN wins a lot during training but benchmarks poorly

Training win rate is measured with exploration active (ε > 0). Benchmark uses pure exploitation (ε = 0). If the bot's greedy policy is weak, more training with a lower epsilon floor is needed. Set **"Epsilon min" to 0.01** and run 5,000 more episodes.

↑ Back to top